



**МИНОБРНАУКИ РОССИИ**  
Федеральное государственное бюджетное образовательное учреждение  
высшего образования  
**«МИРЭА – Российский технологический университет»**  
**РТУ МИРЭА**

---

ИКБ направление «Киберразведка и противодействие угрозам с применением технологий искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной безопасности»

**Практическая работа №6**

по дисциплине «Анализ защищенности систем искусственного интеллекта»

Выполнила:  
Котюкова В.О.

Проверил: Спирин А.А.

Москва 2023

В данной работе мы применим алгоритм для защиты модели. После отработки данного алгоритма мы увидим улучшение в работе модели, которая была подвергнута атаке.

Выполняем импорт необходимых библиотек.

```
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import transforms, datasets
```

Зададим нормализующие преобразования и загрузим набор данных (MNIST), разбить данные на подвыборки

```
transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((0.0,), (1.0,))])
dataset = datasets.MNIST(root = './data', train=True, transform =
transform, download=True)
train_set, val_set = torch.utils.data.random_split(dataset, [50000,
10000])
test_set = datasets.MNIST(root = './data', train=False, transform =
transform, download=True)
train_loader =
torch.utils.data.DataLoader(train_set, batch_size=1, shuffle=True)
val_loader =
torch.utils.data.DataLoader(val_set, batch_size=1, shuffle=True)
test_loader =
torch.utils.data.DataLoader(test_set, batch_size=1, shuffle=True)
print("Training data:", len(train_loader), "Validation
data:", len(val_loader), "Test data: ", len(test_loader))
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
```

```
100%|██████████| 9912422/9912422 [00:00<00:00, 61439365.22it/s]
```

```
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to
./data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubyte.gz
```

```
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
```

```
100%|██████████| 28881/28881 [00:00<00:00, 27872916.20it/s]
```

```
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to
./data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz
```

```
100%|██████████| 1648877/1648877 [00:00<00:00, 41238924.75it/s]
```

```
Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to
./data/MNIST/raw
```

```
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
```

```
100%|██████████| 4542/4542 [00:00<00:00, 9225437.66it/s]
```

```
Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to
./data/MNIST/raw
```

```
Training data: 50000 Validation data: 10000 Test data: 10000
```

Настроим использование графического ускорителя.

```
use_cuda=True
device = torch.device("cuda" if (use_cuda and
torch.cuda.is_available()) else "cpu")
```

Создадим класс НС на основе фреймворка torch.

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)
    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
```

```

x = self.dropout2(x)
x = self.fc2(x)
output = F.log_softmax(x, dim=1)
return output

```

Проверим работоспособность созданного класса НС.

```
model = Net().to(device)
```

Создадим оптимизатор, функцию потерь и трейнер сети.

```

optimizer = optim.Adam(model.parameters(),lr=0.0001, betas=(0.9,
0.999))
criterion = nn.NLLLoss()
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer,
mode='min', factor=0.1, patience=3)

```

Определим функцию обучения сети.

```

def fit(model,device,train_loader,val_loader,epochs):
    data_loader = {'train':train_loader,'val':val_loader}
    print("Fitting the model...")
    train_loss,val_loss=[],[]
    for epoch in range(epochs):
        loss_per_epoch,val_loss_per_epoch=0,0
        for phase in ('train','val'):
            for i,data in enumerate(data_loader[phase]):
                input,label = data[0].to(device),data[1].to(device)
                output = model(input)
                #calculating loss on the output
                loss = criterion(output,label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #update weights
                    optimizer.step()
                    loss_per_epoch+=loss.item()
                else:
                    val_loss_per_epoch+=loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
            print("Epoch: {} Loss: {} Val_Loss:
{}").format(epoch+1,loss_per_epoch/len(train_loader),val_loss_per_epoch
/len(val_loader))
            train_loss.append(loss_per_epoch/len(train_loader))
            val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss,val_loss

```

Обучим модель.

```
loss, val_loss = fit(model, device, train_loader, val_loader, 10)
```

Fitting the model...


```
/usr/local/lib/python3.10/dist-packages/torch/nn/functional.py:1345:  
UserWarning: dropout2d: Received a 2-D input to dropout2d, which is  
deprecated and will result in an error in a future release. To retain  
the behavior and silence this warning, please use dropout instead.  
Note that dropout2d exists to provide channel-wise dropout on inputs  
with 2 spatial dimensions, a channel dimension, and an optional batch  
dimension (i.e. 3D or 4D inputs).
```

```
warnings.warn(warn_msg)
```

```
Epoch: 1 Loss: 0.2806496279286237 Val_Loss: 0.12787441771870753  
Epoch: 2 Loss: 0.11299285361314047 Val_Loss: 0.10075318533363575  
Epoch: 3 Loss: 0.0849774527908649 Val_Loss: 0.09567143414565794  
Epoch: 4 Loss: 0.07608217686837153 Val_Loss: 0.08107877253171125  
Epoch: 5 Loss: 0.0699291657627842 Val_Loss: 0.07728377175676672  
Epoch: 6 Loss: 0.06695735126908428 Val_Loss: 0.0891313714174519  
Epoch: 7 Loss: 0.06077906305478593 Val_Loss: 0.07742812868754662  
Epoch: 8 Loss: 0.06131529724478225 Val_Loss: 0.09425340551983312  
Epoch: 9 Loss: 0.05858486058967601 Val_Loss: 0.08656207511608456  
Epoch: 10 Loss: 0.04640337399357746 Val_Loss: 0.06904277380505101
```

Построим графики потерь при обучении и валидации в зависимости от эпохи.

```
fig = plt.figure(figsize=(5,5))  
plt.plot(np.arange(1,11), loss, "*-", label="Loss")  
plt.plot(np.arange(1,11), val_loss, "o-", label="Val Loss")  
plt.xlabel("Num of epochs")  
plt.legend()  
plt.show()
```



Создадим функции атак FGSM, I-FGSM, MI-FGSM.

```
def fgsm_attack(input,epsilon,data_grad):
    pert_out = input + epsilon*data_grad.sign()
    pert_out = torch.clamp(pert_out, 0, 1)
    return pert_out

def ifgsm_attack(input,epsilon,data_grad):
    pert_out = input + epsilon*data_grad.sign()
    pert_out = torch.clamp(pert_out, 0, 1)
    return pert_out

def mifgsm_attack(input,epsilon,data_grad):
    iter=10
    decay_factor=1.0
    pert_out = input
    alpha = epsilon/iter
    g=0
    for i in range(iter-1):
        g = decay_factor*g + data_grad/torch.norm(data_grad,p=1)
        pert_out = pert_out + alpha*torch.sign(g)
        pert_out = torch.clamp(pert_out, 0, 1)
        if torch.norm((pert_out-input),p=float('inf')) > epsilon:
```

```
        break
    return pert_out
```

Создадим функцию проверки.

```
def test(model, device, test_loader, epsilon, attack):
    correct = 0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data, epsilon, data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data, epsilon, data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data, epsilon, data_grad)
        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
        if (epsilon == 0) and (len(adv_examples) < 5):
            adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
            adv_examples.append( (init_pred.item(), final_pred.item(),
adv_ex) )
        else:
            if len(adv_examples) < 5:
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(),
adv_ex) )
    final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon,
correct, len(test_loader), final_acc))
    return final_acc, adv_examples
```

Построим графики успешности атак (Ассурасу/эпсилон) и примеры выполненных атак в зависимости от степени возмущения epsilon.

```
epsilons = [0, 0.007, 0.01, 0.02, 0.03, 0.05, 0.1, 0.2, 0.3]
for attack in ("fgsm", "ifgsm", "mifgsm"):
    accuracies = []
```

```


examples = []
for eps in epsilons:
    acc, ex = test(model, device, test_loader, eps, attack)
    accuracies.append(acc)
    examples.append(ex)
plt.figure(figsize=(5,5))
plt.plot(epsilons, accuracies, "*-")
plt.title(attack)
plt.xlabel("Epsilon")
plt.ylabel("Accuracy")
plt.show()














































cnt = 0
plt.figure(figsize=(8,10))
for i in range(len(epsilons)):
    for j in range(len(examples[i])):
        cnt += 1
        plt.subplot(len(epsilons), len(examples[0]), cnt)
        plt.xticks([], [])
        plt.yticks([], [])
        if j == 0:
            plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
        orig, adv, ex = examples[i][j]
        plt.title("{} -> {}".format(orig, adv))
        plt.imshow(ex, cmap="gray")
plt.tight_layout()
plt.show()

```


Epsilon: 0	Test Accuracy = 9750 / 10000 = 0.975
Epsilon: 0.007	Test Accuracy = 9722 / 10000 = 0.9722
Epsilon: 0.01	Test Accuracy = 9723 / 10000 = 0.9723
Epsilon: 0.02	Test Accuracy = 9679 / 10000 = 0.9679
Epsilon: 0.03	Test Accuracy = 9613 / 10000 = 0.9613
Epsilon: 0.05	Test Accuracy = 9450 / 10000 = 0.945
Epsilon: 0.1	Test Accuracy = 8837 / 10000 = 0.8837
Epsilon: 0.2	Test Accuracy = 6246 / 10000 = 0.6246
Epsilon: 0.3	Test Accuracy = 3601 / 10000 = 0.3601



















































Eps: 0	1 -> 1 	9 -> 9 	1 -> 1 	7 -> 7 	7 -> 7 
Eps: 0.007	1 -> 1 	4 -> 4 	9 -> 9 	0 -> 0 	8 -> 8 
Eps: 0.01	7 -> 7 	9 -> 9 	9 -> 9 	2 -> 2 	3 -> 3 
Eps: 0.02	0 -> 0 	5 -> 5 	8 -> 8 	5 -> 5 	8 -> 8 
Eps: 0.03	8 -> 8 	4 -> 4 	2 -> 2 	0 -> 0 	0 -> 0 
Eps: 0.05	5 -> 5 	1 -> 1 	2 -> 2 	1 -> 1 	7 -> 7 
Eps: 0.1	5 -> 5 	9 -> 9 	8 -> 8 	8 -> 8 	2 -> 2 
Eps: 0.2	8 -> 8 	3 -> 3 	6 -> 6 	3 -> 3 	0 -> 0 
Eps: 0.3	0 -> 2 	4 -> 4 	1 -> 1 	8 -> 8 	9 -> 4 














































Epsilon: 0 Test Accuracy = 9751 / 10000 = 0.9751  
Epsilon: 0.007 Test Accuracy = 9715 / 10000 = 0.9715  
Epsilon: 0.01 Test Accuracy = 9709 / 10000 = 0.9709  
Epsilon: 0.02 Test Accuracy = 9688 / 10000 = 0.9688  
Epsilon: 0.03 Test Accuracy = 9612 / 10000 = 0.9612  
Epsilon: 0.05 Test Accuracy = 9476 / 10000 = 0.9476  
Epsilon: 0.1 Test Accuracy = 8841 / 10000 = 0.8841  
Epsilon: 0.2 Test Accuracy = 6273 / 10000 = 0.6273  
Epsilon: 0.3 Test Accuracy = 3484 / 10000 = 0.3484



Eps: 0	0 -> 0 	9 -> 9 	2 -> 2 	2 -> 2 	1 -> 1 
Eps: 0.007	4 -> 4 	4 -> 4 	8 -> 8 	7 -> 7 	9 -> 9 
Eps: 0.01	3 -> 3 	9 -> 9 	8 -> 8 	5 -> 5 	1 -> 1 
Eps: 0.02	7 -> 7 	3 -> 3 	9 -> 9 	0 -> 0 	4 -> 4 
Eps: 0.03	5 -> 5 	7 -> 7 	5 -> 5 	4 -> 4 	9 -> 9 
Eps: 0.05	0 -> 0 	3 -> 3 	8 -> 8 	9 -> 9 	8 -> 8 
Eps: 0.1	0 -> 9 	6 -> 6 	1 -> 1 	9 -> 9 	2 -> 2 
Eps: 0.2	1 -> 2 	2 -> 7 	1 -> 1 	1 -> 1 	7 -> 7 
Eps: 0.3	2 -> 2 	7 -> 3 	2 -> 8 	1 -> 2 	6 -> 8 

Epsilon: 0 Test Accuracy = 9738 / 10000 = 0.9738  
Epsilon: 0.007 Test Accuracy = 9728 / 10000 = 0.9728  
Epsilon: 0.01 Test Accuracy = 9713 / 10000 = 0.9713  
Epsilon: 0.02 Test Accuracy = 9687 / 10000 = 0.9687  
Epsilon: 0.03 Test Accuracy = 9634 / 10000 = 0.9634  
Epsilon: 0.05 Test Accuracy = 9527 / 10000 = 0.9527  
Epsilon: 0.1 Test Accuracy = 8986 / 10000 = 0.8986  
Epsilon: 0.2 Test Accuracy = 6933 / 10000 = 0.6933  
Epsilon: 0.3 Test Accuracy = 4250 / 10000 = 0.425



Eps: 0	6 -> 6 	3 -> 3 	2 -> 2 	5 -> 5 	4 -> 4 
Eps: 0.007	3 -> 3 	0 -> 0 	0 -> 0 	6 -> 6 	6 -> 6 
Eps: 0.01	6 -> 6 	7 -> 7 	8 -> 8 	1 -> 1 	3 -> 3 
Eps: 0.02	5 -> 5 	2 -> 2 	5 -> 5 	1 -> 1 	0 -> 0 
Eps: 0.03	4 -> 4 	9 -> 9 	7 -> 7 	6 -> 6 	5 -> 5 
Eps: 0.05	2 -> 2 	1 -> 1 	3 -> 3 	6 -> 6 	8 -> 8 
Eps: 0.1	9 -> 4 	0 -> 0 	7 -> 7 	5 -> 5 	8 -> 8 
Eps: 0.2	5 -> 5 	6 -> 0 	3 -> 3 	5 -> 5 	1 -> 1 
Eps: 0.3	5 -> 3 	0 -> 5 	3 -> 1 	5 -> 8 	0 -> 5 

Создадим 2 класса НС.

```
class NetF(nn.Module):
    def __init__(self):
        super(NetF, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x

class NetF1(nn.Module):
    def __init__(self):
        super(NetF1, self).__init__()
        self.conv1 = nn.Conv2d(1, 16, 3, 1)
        self.conv2 = nn.Conv2d(16, 32, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(4608, 64)
        self.fc2 = nn.Linear(64, 10)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        return x
```

Переопределим функцию обучения и тестирования.

```
def
fit(model,device,optimizer,scheduler,criterion,train_loader,val_loader
,Temp,epochs):
    data_loader = {'train':train_loader,'val':val_loader}
    print("Fitting the model...")
    train_loss,val_loss=[],[]
    for epoch in range(epochs):
        loss_per_epoch,val_loss_per_epoch=0,0
        for phase in ('train','val'):
            for i,data in enumerate(data_loader[phase]):
                input,label = data[0].to(device),data[1].to(device)
                output = model(input)
                output = F.log_softmax(output/Temp,dim=1)
                #calculating loss on the output
                loss = criterion(output,label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #update weights
                    optimizer.step()
                    loss_per_epoch+=loss.item()
                else:
                    val_loss_per_epoch+=loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
            print("Epoch: {} Loss: {} Val_Loss:
{}".format(epoch+1,loss_per_epoch/len(train_loader),val_loss_per_epoch
/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss,val_loss
def test(model,device,test_loader,epsilon,Temp,attack):
    correct=0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        output = F.log_softmax(output/Temp,dim=1)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data,epsilon,data_grad)
```



```

elif attack == "ifgsm":
    perturbed_data = ifgsm_attack(data,epsilon,data_grad)
elif attack == "mifgsm":
    perturbed_data = mifgsm_attack(data,epsilon,data_grad)
output = model(perturbed_data)
final_pred = output.max(1, keepdim=True)[1]
if final_pred.item() == target.item():
    correct += 1
    if (epsilon == 0) and (len(adv_examples) < 5):
        adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
        adv_examples.append( (init_pred.item(), final_pred.item(),
adv_ex) )
    else:
        if len(adv_examples) < 5:
            adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
            adv_examples.append( (init_pred.item(), final_pred.item(),
adv_ex) )
    final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon,
correct, len(test_loader), final_acc))
    return final_acc,adv_examples

```

Создадим функцию защиты методом дистилляции.

```

def
defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilon
s):
    modelF = NetF().to(device)
    optimizerF = optim.Adam(modelF.parameters(),lr=0.0001, betas=(0.9,
0.999))
    schedulerF = optim.lr_scheduler.ReduceLROnPlateau(optimizerF,
mode='min', factor=0.1, patience=3)
    modelF1 = NetF1().to(device)
    optimizerF1 = optim.Adam(modelF1.parameters(),lr=0.0001, betas=(0.9,
0.999))
    schedulerF1 = optim.lr_scheduler.ReduceLROnPlateau(optimizerF1,
mode='min', factor=0.1, patience=3)
    criterion = nn.NLLLoss()

lossF,val_lossF=fit(modelF,device,optimizerF,schedulerF,criterion,train_loader,val_loader,Temp,epochs)
fig = plt.figure(figsize=(5,5))
plt.plot(np.arange(1,epochs+1), lossF, "*-",label="Loss")
plt.plot(np.arange(1,epochs+1), val_lossF,"o-",label="Val Loss")
plt.title("Network F")
plt.xlabel("Num of epochs")
plt.legend()
plt.show()
#converting target labels to soft labels

```

```

for data in train_loader:
    input, label = data[0].to(device), data[1].to(device)
    softlabel = F.log_softmax(modelF(input), dim=1)
    data[1] = softlabel

lossF1, val_lossF1 = fit(modelF1, device, optimizerF1, schedulerF1, criterion,
                        train_loader, val_loader, Temp, epochs)
fig = plt.figure(figsize=(5, 5))
plt.plot(np.arange(1, epochs+1), lossF1, "*-", label="Loss")
plt.plot(np.arange(1, epochs+1), val_lossF1, "o-", label="Val Loss")
plt.title("Network F'")
plt.xlabel("Num of epochs")
plt.legend()
plt.show()
model = NetF1().to(device)
model.load_state_dict(modelF1.state_dict())
for attack in ("fgsm", "ifgsm", "mifgsm"):
    accuracies = []
    examples = []
    for eps in epsilons:
        acc, ex = test(model, device, test_loader, eps, attack)
        accuracies.append(acc)
        examples.append(ex)
    plt.figure(figsize=(5, 5))
    plt.plot(epsilons, accuracies, "*-")
    plt.title(attack)
    plt.xlabel("Epsilon")
    plt.ylabel("Accuracy")
    plt.show()

cnt = 0
plt.figure(figsize=(8, 10))
for i in range(len(epsilons)):
    for j in range(len(examples[i])):
        cnt += 1
        plt.subplot(len(epsilons), len(examples[0]), cnt)
        plt.xticks([], [])
        plt.yticks([], [])
        if j == 0:
            plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
        orig, adv, ex = examples[i][j]
        plt.title("{} -> {}".format(orig, adv))
        plt.imshow(ex, cmap="gray")
plt.tight_layout()
plt.show()


```

Получим результаты оценки защищенных сетей.

```
Temp=100
epochs=10
epsilons=[0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons)
```

Fitting the model...

```
Epoch: 1 Loss: 0.5662047301409073 Val_Loss: 0.00018231202587485313
Epoch: 2 Loss: 0.3385203243374718 Val_Loss: 1.2374566635116935e-06
Epoch: 3 Loss: 0.25428086451666876 Val_Loss: 3.757412723643938e-07
Epoch: 4 Loss: 0.19064711203381343 Val_Loss: 0.00048478846549874106
Epoch: 5 Loss: 0.15228855001899314 Val_Loss: 1.921464577317238e-05
Epoch: 6 Loss: 0.1285232472288784 Val_Loss: 1.4979687286540866e-06
Epoch: 7 Loss: 0.11097932044940916 Val_Loss: 2.742809010669589e-08
Epoch: 8 Loss: 0.10269630508215549 Val_Loss: 7.760227151720756e-09
Epoch: 9 Loss: 0.09628898256426938 Val_Loss: 1.327698947279714e-05
Epoch: 10 Loss: 0.09052555274544327 Val_Loss: 1.2294865399599076e-05
```




Fitting the model...

```
Epoch: 1 Loss: 0.6954153548508452 Val_Loss: 3.698446452617645e-05
Epoch: 2 Loss: 0.45832563545066807 Val_Loss: 1.771830189973116e-05
Epoch: 3 Loss: 0.3865080458324801 Val_Loss: 1.3489030115306378e-05
```

```

Epoch: 4 Loss: 0.3318358400624503 Val_Loss: 5.6694397889077666e-05
Epoch: 5 Loss: 0.2763107574661768 Val_Loss: 0.00014575666207820178
Epoch: 6 Loss: 0.2445015275636196 Val_Loss: 4.863006553932791e-07
Epoch: 7 Loss: 0.22031390263005912 Val_Loss: 1.8023547500524728e-05
Epoch: 8 Loss: 0.20348487243809338 Val_Loss: 3.60193886276912e-07
Epoch: 9 Loss: 0.19343684324045543 Val_Loss: 5.017868280410767e-05
Epoch: 10 Loss: 0.1826015582433685 Val_Loss: 0.00013195873200893402

```

















































```

Epsilon: 0 Test Accuracy = 9263 / 10000 = 0.9263
Epsilon: 0.007 Test Accuracy = 9170 / 10000 = 0.917
Epsilon: 0.01 Test Accuracy = 9107 / 10000 = 0.9107
Epsilon: 0.02 Test Accuracy = 9016 / 10000 = 0.9016
Epsilon: 0.03 Test Accuracy = 8894 / 10000 = 0.8894
Epsilon: 0.05 Test Accuracy = 8496 / 10000 = 0.8496
Epsilon: 0.1 Test Accuracy = 7037 / 10000 = 0.7037
Epsilon: 0.2 Test Accuracy = 2894 / 10000 = 0.2894
Epsilon: 0.3 Test Accuracy = 1105 / 10000 = 0.1105
Epsilon: 0 Test Accuracy = 9214 / 10000 = 0.9214
Epsilon: 0.007 Test Accuracy = 9185 / 10000 = 0.9185
Epsilon: 0.01 Test Accuracy = 9117 / 10000 = 0.9117
Epsilon: 0.02 Test Accuracy = 9012 / 10000 = 0.9012

```

```
Epsilon: 0.03    Test Accuracy = 8906 / 10000 = 0.8906
Epsilon: 0.05    Test Accuracy = 8507 / 10000 = 0.8507
Epsilon: 0.1     Test Accuracy = 6995 / 10000 = 0.6995
Epsilon: 0.2     Test Accuracy = 2945 / 10000 = 0.2945
Epsilon: 0.3     Test Accuracy = 1156 / 10000 = 0.1156
Epsilon: 0 Test Accuracy = 9203 / 10000 = 0.9203
Epsilon: 0.007   Test Accuracy = 9167 / 10000 = 0.9167
Epsilon: 0.01    Test Accuracy = 9106 / 10000 = 0.9106
Epsilon: 0.02    Test Accuracy = 9008 / 10000 = 0.9008
Epsilon: 0.03    Test Accuracy = 8911 / 10000 = 0.8911
Epsilon: 0.05    Test Accuracy = 8476 / 10000 = 0.8476
Epsilon: 0.1     Test Accuracy = 7052 / 10000 = 0.7052
Epsilon: 0.2     Test Accuracy = 2897 / 10000 = 0.2897
Epsilon: 0.3     Test Accuracy = 1136 / 10000 = 0.1136
```



Eps: 0	7 -> 7 	0 -> 0 	2 -> 2 	7 -> 7 	1 -> 1 
Eps: 0.007	0 -> 0 	6 -> 6 	7 -> 7 	9 -> 9 	0 -> 0 
Eps: 0.01	2 -> 2 	9 -> 9 	9 -> 9 	1 -> 1 	9 -> 9 
Eps: 0.02	6 -> 6 	7 -> 7 	2 -> 2 	4 -> 4 	3 -> 3 
Eps: 0.03	1 -> 1 	8 -> 8 	8 -> 8 	7 -> 7 	6 -> 6 
Eps: 0.05	8 -> 8 	0 -> 0 	6 -> 6 	2 -> 1 	7 -> 3 
Eps: 0.1	8 -> 8 	4 -> 4 	9 -> 7 	8 -> 1 	7 -> 7 
Eps: 0.2	4 -> 9 	5 -> 8 	2 -> 0 	9 -> 8 	9 -> 8 
Eps: 0.3	3 -> 8 	7 -> 4 	1 -> 2 	8 -> 3 	3 -> 0 

Дистилляция (или обучение с учителем) - это метод, при котором более крупная модель (учитель) используется для обучения более компактной модели (ученика) путем передачи знаний из учителя в ученика. В контексте защиты от адверсариальных атак, дистилляция может применяться для создания модели, более устойчивой к таким атакам.

Учительская модель может обучить ученическую модель выделять более стабильные и робастные паттерны в данных, что может помочь предотвратить некоторые виды адверсариальных атак. Дистилляция может сделать модель менее чувствительной к мелким изменениям входных данных, затруднить успешное проведение атаки, но также может снизить точность модели (особенно если не подобрать оптимальные параметры процесса дистилляции), повлечь за собой увеличение вычислительной сложности обучения и инференса.

В нашем случае, мы можем видеть такие показатели точности атак:

```
fgsm    - 0.97 / 0.36  
ifgsm   - 0.97 / 0.34  
mifgsm  - 0.97 / 0.42
```

А точность модели остается в пределах 0.9